

Project Stage IV

Merging tables

Team

Anna Chang
Sripradha Karkala
Simmi Pateriya

Data Description

We used the songs and tracks tables from the movie-track-song dataset that was provided in the previous project stage. Each tuple in these tables describe a song entity. The original tables contained 961,593 tuples and 734,485 tuples for songs and tracks respectively. We reduced the size to 17,731 (*sample_A*) and 10,000 (*sample_B*) after downsampling.

Merging process

The downsampled datasets (*sample_A* and *sample_B*) and a matcher table (*matches.csv*) files were generated as a result of the previous stage on data matching.

Schema of *sample_A*

id	title	artist_name	year
457	Mango Tree	Angus & Julia Stone	2006

Schema of *sample_B*

id	title	year	episode	song	artists
140954	Lockie Leonard	2007	New and Improved (#2.1)	Mango Tree	angus stone+julia stone+angus

Schema of the matcher table (*matches*)

<i>ltable_id</i>	<i>rtable_id</i>	<i>Ltable_title</i>	<i>ltable_artist_name</i>	<i>ltable_year</i>	<i>rtable_title</i>	<i>rtable_year</i>	<i>rtable_episode</i>	<i>rtable_song</i>	<i>rtable_artists</i>	<i>Lin. reg</i>	<i>Log reg</i>	<i>svm</i>
457	140954	Mango Tree	Angus & Julia Stone	2006	Lockie Leonard	2007	New and Improved (#2.1)	Mango Tree	angus stone + julia stone + angus	1	1	1

Using the *ltable_id* and *rtable_id* from the matcher table, we merged the data from the downsized samples to generate a single merged table (*merged.csv*). The matcher table was manually cleaned to exclude any false positive and include any true negatives. These rows are highlighted in [Matches_Cleaned.xlsx](#)

From the data matching stage we know that *artists_name* and *artists* from *sample_A* and *sample_B* respectively refer to the same real world entity. Similarly, *title* from *sample_A* and *song* from *sample_B* refer to the same real world entity. We pick the matcher for Logistic Regression as it has better accuracy.

The following rules were used for generating the merged table.

Column	Table	Rule
<i>id</i>	<i>Sample_A</i>	<i>song_id</i>
<i>id</i>	<i>Sample_B</i>	<i>track_id</i>
<i>title</i>	<i>Sample_B</i>	<i>title</i>
<i>artist_name, artists</i>	<i>Sample_A, Sample_B</i>	Take the longer string, which includes more information on the artists. Since this is based on string length, missing values are taken care of. This column is stored as <i>artists</i>
<i>title, song</i>	<i>Sample_A, Sample_B</i>	Take the longer string, which includes more information about the song. Since this is based on the string length, missing values are taken care of. This column is stored as <i>song</i> .

year	Sample_A	year
episode	Sample_B	episode
year	Sample_B	episode_year

The merged table contains 1331 tuples. And the final schema of the merged tables looks as follows:

song_id	track_id	title	artists	song	year	episode	episode_year
457	140954	Lockie Leonard	angus stone+julia stone+angus	Mango Tree	2006	New and Improved (#2.1)	2007
1203	444229	Frank Sinatra: A Man and His Music + Ella + Jobim	paul mann+stefan weierstraß+frank sinatra	Put Your Dreams Away	1958		1967
1935	121170	Intervention	scott klass+the davenports	Five Steps	2000	Adam (#9.2)	2005
1935	121498	Intervention	scott klass+the davenports	Five Steps	2000	Salina and Troy (#2.5)	2005
4202	541635	Ministry: Tapes of Wrath	ministry	The Land of Rape & Honey (Live)	0		2000
4432	208737	Shameless	the high strung	The Luck You Got	2005	The Sins of My Caretaker (#3.5)	2011

Source code

```
import pandas as pd

'''Rules:
save id from songs table as song_id
save id from tracks table as track_id
save title from tracks table
save artists names from whichever is long
save song from tracks if there else title
from songs
save year from songs
save episode from tracks
save episode year from tracks
'''

def apply_rules(tup1, tup2):
    song_id = tup1['id'].iloc[0]
    track_id = tup2['id'].iloc[0]
    title = tup2['title'].iloc[0]
    tup2_artist = tup2['artists'].iloc[0]
    tup1_artist = tup1['artist_name'].iloc[0]
    if len(tup2_artist) >= len(tup1_artist):
        artists = tup2_artist
    else:
        artists = tup1_artist
    tup2_song = tup2['song'].iloc[0]
    tup1_song = tup1['title'].iloc[0]
    if len(str(tup2_song)) >= len(tup1_song):
        song = tup2_song
    else:
        song = tup1_song
    year = tup1['year'].iloc[0]
    episode = tup2['episode'].iloc[0]
    episode_year = tup2['year'].iloc[0]

    data = [{'song_id':song_id, 'track_id':track_id, 'title':title, \
            'artists':artists, 'song':song , \
            'year':year, 'episode':episode, 'episode_year':episode_year}]
    return data

def save_to_csv(merged):
    merged.to_csv('./merged.csv', index=False)

'''create a schema as union of both tables
For the true matches we found in Linear regression
find corresponding tables from songs and tracks table.
Apply rules to combine the common attributes.
Write the final merged table as merged.csv
```

```

'''
def merge_table(table1, table2, matcher):
    schema = ['song_id', 'track_id', 'title', 'artists', 'song', 'year',\
              'episode', 'episode_year']
    merged_table = pd.DataFrame(columns=schema)
    for index, match in matcher.iterrows():
        if match['logistic regression'] == 1:
            tup_table1 = table1.loc[table1.id == match['ltable_id']]
            tup_table2 = table2.loc[table2.id == match['rtable_id']]
            data = apply_rules(tup_table1, tup_table2)
            df = pd.DataFrame(data, columns=schema)
            merged_table = merged_table.append(df)
    save_to_csv(merged_table)

if __name__ == '__main__':
    '''Read downsampled files and matcher
    sampleA - downsampled song dataset
    sampleB - downsampled tracks dataset
    Matches - Our matcher's output
    '''
    table1 = pd.read_csv('./sampleA.csv', index_col=False)
    #if pandas version more than 0.17 use table1.sort_values(by=['id'], ascending=True)
    table1.sort('id', ascending=True, inplace=True)
    table2 = pd.read_csv('./sampleB.csv', index_col=False)
    table2.sort('id', ascending=True, inplace=True)
    matcher = pd.read_csv('./Matches_Cleaned.csv', index_col=False)
    matcher.sort('ltable_id', ascending=True, inplace=True)
    merge_table(table1, table2, matcher)

```